

Surya Group OF Institution
Naan Muthalvan
IBM

IBM - Artificial Intelligence
VIMAL RAJ M
422221104047
Team - 05

Sentiment analysis for marketing

Sentiment analysis is a technique that uses natural language processing and machine learning to analyze the emotions and opinions expressed in text data. It can help marketers to understand how their customers feel about their products, services, or brand, and to improve their marketing strategies accordingly.

There are many ways to perform sentiment analysis using Python, one of the most popular programming languages for data science.

Some of the possible methods are:

- Using pre-trained models from Hugging Face, a platform that provides a large collection of state-of-the-art natural language processing models. You can use the `pipeline` function from the `transformers` library to create a sentiment analysis pipeline that can take any text input and return a label (positive or negative) and a score (confidence level).

```
pip install -q transformers
from transformers import pipeline
sentiment_pipeline = pipeline("sentiment-analysis")
data = ["I love you", "I hate you"]
sentiment_pipeline(data)
```

- Using the `TextBlob` library, which is a simple and intuitive way to perform sentiment analysis in Python. It has a built-in sentiment analyzer that returns a polarity (between -1 and 1) and a subjectivity (between 0 and 1) score for any text input.

```
pip install -q textblob
from textblob import TextBlob
data = ["I love you", "I hate you"]
for text in data:
    blob = TextBlob(text)
    print(blob.sentiment)
```

- Using the `NLTK` library, which is a comprehensive toolkit for natural language processing in Python. It has various modules and resources for sentiment analysis, such as the `VADER` (Valence Aware Dictionary and sEntiment Reasoner) tool, which is a lexicon-based approach that uses a list of words with predefined sentiment scores to calculate the overall sentiment of a text input.

```
pip install -q nltk
import nltk
nltk.download('vader_lexicon')
from nltk.sentiment.vader import SentimentIntensityAnalyzer
sia = SentimentIntensityAnalyzer()
data = ["I love you", "I hate you"]
for text in data:
    print(sia.polarity_scores(text))
```

- Building your own sentiment analysis model using TensorFlow, which is an open-source framework for machine learning and deep learning. You can use the keras API to create a neural network model that can learn from labeled text data and predict the sentiment of new text inputs.

Table of Contents:

- [How to Perform Sentiment Analysis in Python?](#)
 1. [Python Pre-Requisites](#)
 2. [Reading the Dataset](#)
 3. [Data Preprocessing](#)
 4. [TF-IDF Transformation](#)
 5. [Building and Evaluating the ML Model](#)
- [Sentiment Analysis with Python: Next Steps](#)

How to Perform Sentiment Analysis in Python?

You're probably already familiar with Python, but if not – it is a powerful programming language with an intuitive syntax. Not to mention it's one of the most popular choices across the data science community, which makes it perfect for our tutorial.

Step 1: Python Pre-Requisites :

First things first: installing the necessary equipment. You need a Python IDE – I suggest using Jupyter. (If you don't already have it, follow this [Jupyter Notebook tutorial](#) to set it up on your device.)

Make sure to have the following libraries installed as well: NumPy, pandas, Matplotlib, seaborn, Regex, and scikit-learn.

Step 2: Reading the Dataset :

Let's start by loading the dataset into Python and reading the head of the data frame:

```
import pandas as pd
df = pd.read_csv('tripadvisor_hotel_reviews.csv')
df.head()
```

	Review	Rating
0	nice hotel expensive parking got good deal sta...	4
1	ok nothing special charge diamond member hilto...	2
2	nice rooms not 4* experience hotel monaco seat...	3
3	unique, great stay, wonderful time hotel monac...	5
4	great stay great stay, went seahawk game aweso...	5

```
len(df.index) # 20491
```

Step 3: Data Preprocessing :

As we already know the TripAdvisor dataset has 2 variables – user reviews and ratings, which range from 1 to 5. We will use “Ratings” to create a new variable called “Sentiment.” In it, we will add 2 categories of sentiment as follows:

- 0 to 1 will be encoded as -1 as they indicate negative sentiment
- 3 will be labeled as 0 as it has a neutral sentiment
- 4 and 5 will be labeled as +1 as they indicate positive sentiment

```
import numpy as np
def create_sentiment(rating):
    if rating==1 or rating==2:
        return -1 # negative sentiment
    elif rating==4 or rating==5:
        return 1 # positive sentiment
    else:
        return 0 # neutral sentiment
df['Sentiment'] = df['Rating'].apply(create_sentiment)
```

	Review	Rating	Sentiment
0	nice hotel expensive parking got good deal sta...	4	1
1	ok nothing special charge diamond member hilt...	2	-1
2	nice rooms not experience hotel monaco seattl...	3	0
3	unique great stay wonderful time hotel monaco ...	5	1
4	great stay great stay went seahawk game awesom...	5	1

```
from sklearn.feature_extraction.text import re
def clean_data(review):
    no_punc = re.sub(r'^\w\s]', '', review)
    no_digits = ''.join([i for i in no_punc if not i.isdigit()])
    return(no_digits)
```

```
df['Review'] = df['Review'].apply(clean_data)
df['Review'][0]
```

Step 4: TF-IDF Transformation :

Rather than focusing on individual pieces, inverse document frequency measures how many times a word is repeated across a set of documents. And opposite of the previous metric, here the higher frequency is – the lower the relevance. This helps the algorithm eliminate naturally occurring words such as “a”, “the”, “and”, etc, as they will appear frequently across all documents in a corpus.

Now that you understand how TF-IDF works, let’s use this algorithm to vectorize our data:

```
from sklearn.feature_extraction.text import TfidfVectorizer
tfidf = TfidfVectorizer(strip_accents=None,
lowercase=False,
preprocessor=None)
X = tfidf.fit_transform(df['Review'])
```

Step 5: Building and Evaluating the Machine Learning Model :

We can now train our algorithm on the review data to classify its sentiment into 3 categories:

- *Positive*
- *Negative*
- *Neutral*

```
from sklearn.model_selection import train_test_split
y = df['Sentiment'] # target variable
X_train, X_test, y_train, y_test = train_test_split(X,y)
```

```
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression(solver='liblinear')
lr.fit(X_train,y_train) # fit the model
preds = lr.predict(X_test) # make predictions
```

Finally, evaluate the performance:

```
from sklearn.metrics import accuracy_score accuracy_score
(preds,y_test) # 0.86
```

Our model has an accuracy of approximately 0.86, which is quite good.

And that concludes our tutorial! For a better understanding of the concept, here is the complete sentiment analysis Python code I've used:

```
import pandas as pd
import numpy as np
from sklearn.feature_extraction.text import re
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
df = pd.read_csv('tripadvisor_hotel_reviews.csv')
def create_sentiment(rating):

res = 0 # neutral sentiment

if rating==1 or rating==2:
res = -1 # negative sentiment
elif rating==4 or rating==5:
```

```
res = 1 # positive sentiment

return res

df['Sentiment'] = df['Rating'].apply(create_sentiment)

def clean_data(review):

    no_punc = re.sub(r'^\w\s', '', review)
    no_digits = "".join([i for i in no_punc if not i.isdigit()])

    return(no_digits)

df['Review'] = df['Review'].apply(clean_data)

tfidf = TfidfVectorizer(strip_accents=None,
lowercase=False,
preprocessor=None)
X = tfidf.fit_transform(df['Review'])
y = df['Sentiment']
X_train, X_test, y_train, y_test = train_test_split(X,y)
lr = LogisticRegression(solver='liblinear')
lr.fit(X_train,y_train)
preds = lr.predict(X_test)
accuracy_score(preds,y_test)
```